Statistical Error Metrics & Model Evaluation

Margin of Error, MSE, and R-squared Explained

Zamir Hussain
Department of Mathematics
University of Wah

February 11, 2026

Abstract

These comprehensive lecture notes cover three fundamental error metrics in statistics and machine learning: Margin of Error (MOE), Mean Squared Error (MSE), and R-squared (R^2). Each concept is explained with clear definitions, mathematical formulations, step-by-step examples, and practical applications. Designed for students and practitioners, these notes bridge theoretical foundations with real-world applications in data science.

Contents

1 Introduction to Statistical Error Metrics

Learning Objectives

Learning Objectives:

- Understand the role of error metrics in statistical inference and machine learning
- Differentiate between sampling error and prediction error
- Recognize when to use different error metrics based on context
- Apply error metrics to evaluate model performance and estimate precision

Week 1: Foundations

Why Measure Errors?

- **Quantify Uncertainty:** How confident are we in our estimates?
- **Model Evaluation:** How well does our model predict?
- **Decision Making:** When to trust predictions and when to be cautious
- **Improvement:** Identify areas for model enhancement

Theoretical Foundations

Fundamental Concepts:

- **Error:** Difference between observed and expected values
- **Bias vs Variance:** Trade-off in model complexity
- **Sampling Error:** Due to observing a sample rather than entire population
- **Prediction Error:** Due to model limitations and noise in data

2 Margin of Error (MOE)

2.1 Definition and Conceptual Understanding

The **Margin of Error (MOE)** quantifies the uncertainty in a statistical estimate. It represents the maximum expected difference between the sample estimate and the true population parameter, with a specified level of confidence.

Theoretical Foundations

Key Characteristics:

- Used primarily in **confidence intervals**
- Measures **sampling error**, not measurement error
- Depends on: sample size, population variability, confidence level
- Expressed as ± value around point estimate

2.2 Mathematical Formulation

For Population Mean (Large Samples):

$$MOE = Z_{\alpha/2} \times \frac{s}{\sqrt{n}}$$

Where:

- $Z_{\alpha/2}$ = Critical value from standard normal distribution
- s = Sample standard deviation
- n = Sample size

For Proportions:

$$MOE = Z_{\alpha/2} \times \sqrt{\frac{p(1-p)}{n}}$$

Where p is the sample proportion.

2.3 Common Critical Values

Confidence Level	α	$Z_{\alpha/2}$
90%	0.10	1.645
95%	0.05	1.960
99%	0.01	2.576

2.4 Step-by-Step Example

Problem: Calculate the Margin of Error for a survey with:

- Confidence Level: 99%
- Sample Standard Deviation: $s = 23.44$
- Sample Size: $n = 32$

Solution:
Step 1: Identify critical Z-value For 99% confidence, $Z_{\alpha/2} = 2.576$
Step 2: Calculate standard error

$$SE = \frac{s}{\sqrt{n}} = \frac{23.44}{\sqrt{32}} = \frac{23.44}{5.6569} = 4.142$$

Step 3: Calculate MOE

$$MOE = Z_{\alpha/2} \times SE = 2.576 \times 4.142 = 10.67$$

Interpretation: We are 99% confident that our sample estimate is within ±10.67 units of the true population value.

2.5 Sample Size Determination

Required Sample Size:

$$n = \left(\frac{Z_{\alpha/2} \times s}{MOE} \right)^2$$

Key Insight: To **halve** the MOE, you need to **quadruple** the sample size.

3 Mean Squared Error (MSE)

3.1 Definition and Conceptual Understanding

Mean Squared Error (MSE) measures the average squared difference between predicted values and actual values. It quantifies the quality of an estimator or predictor.

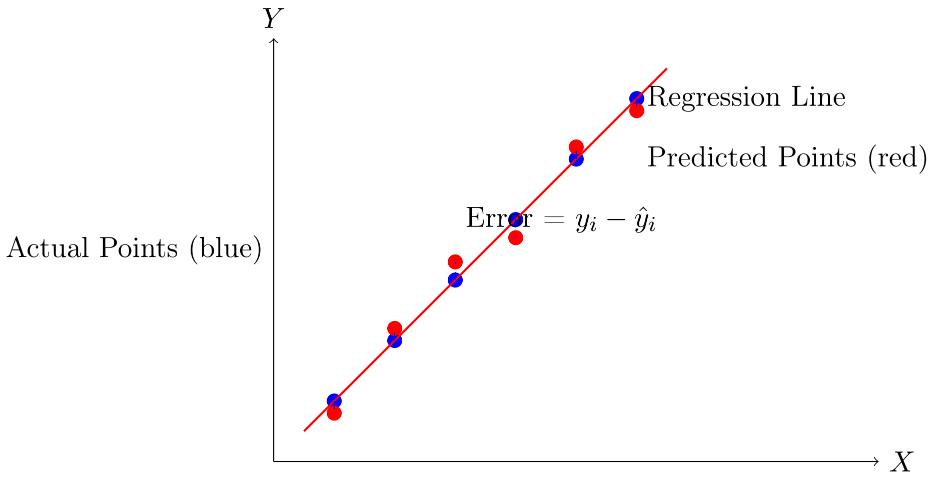

3.2 Mathematical Formulation

Mean Squared Error:

$$MSE = \frac{1}{n} \sum_{i=1}^{n} (y_i - \hat{y}_i)^2$$

Where:

- y_i = Actual value for observation i
- $\hat{y}_i$ = Predicted value for observation i
- n = Number of observations

3.3 Why Square the Errors?

Theoretical Foundations

Advantages of Squaring:

- **Emphasizes Larger Errors:** Squares magnify larger discrepancies
- **Always Positive:** Eliminates cancellation of positive and negative errors
- **Mathematical Convenience:** Differentiable everywhere
- **Statistical Properties:** Related to variance and optimal for Gaussian errors

3.4 Root Mean Squared Error (RMSE)

Root Mean Squared Error:

$$RMSE = \sqrt{MSE} = \sqrt{\frac{1}{n}\sum_{i=1}^{n}(y_i - \hat{y}_i)^2}$$

Advantage: RMSE is in the same units as the original data, making interpretation easier.

3.5 Properties of MSE

- **Range:** $0 \leq MSE < \infty$
- **Perfect Model:** $MSE = 0$ when all predictions are perfect
- **Units:** Square of the original units
- **Sensitivity:** Highly sensitive to outliers due to squaring

4 R-squared (Coefficient of Determination)

4.1 Definition and Conceptual Understanding

R-squared (R^2) or the **coefficient of determination** measures the proportion of variance in the dependent variable that is predictable from the independent variable(s). It quantifies the goodness of fit of a regression model.

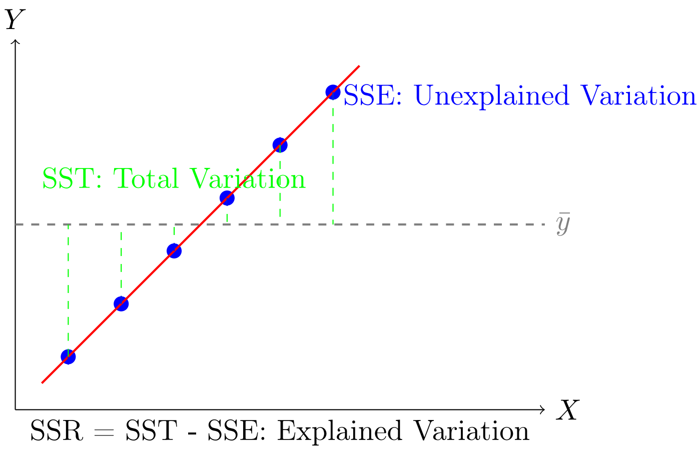

4.2 Mathematical Formulation

R-squared Formula:

$$R^2 = 1 - \frac{SSE}{SST} = 1 - \frac{\sum_{i=1}^{n}(y_i - \hat{y}_i)^2}{\sum_{i=1}^{n}(y_i - \bar{y})^2}$$

Where:

- SSE = Sum of Squared Errors (Unexplained variation)
- SST = Total Sum of Squares (Total variation)
- $\bar{y}$ = Mean of actual values

Alternative Formulation:

$$R^2 = \frac{SSR}{SST} = \frac{\sum_{i=1}^{n}(\hat{y}_i - \bar{y})^2}{\sum_{i=1}^{n}(y_i - \bar{y})^2}$$

Where SSR = Sum of Squares due to Regression (Explained variation).

4.3 Interpretation Guide

R² Value	Quality	Interpretation
0.9 - 1.0	Excellent	Model explains 90-100% of variance. Strong predictive power.
0.7 - 0.9	Good	Model explains 70-90% of variance. Reliable predictions.
0.5 - 0.7	Moderate	Model explains 50-70% of variance. Acceptable for some applications.
0.3 - 0.5	Weak	Model explains 30-50% of variance. Limited predictive value.
0.0 - 0.3	Poor	Model explains less than 30% of variance. Little explanatory power.
Negative	Very Poor	Model performs worse than predicting the mean.

4.4 Step-by-Step Numerical Example

Problem: Calculate R² for the following data:

Actual (y_i)	Predicted ($\hat{y}_i$)
4	2
2	3
3	3.5
5	4
1	4.5

Solution Steps:
Step 1: Calculate mean of actual values

$$\bar{y} = \frac{4+2+3+5+1}{5} = \frac{15}{5} = 3$$

Step 2: Calculate SSE (Sum of Squared Errors)

$$\begin{aligned}
SSE &= \sum (y_i - \hat{y}_i)^2 \\
&= (4-2)^2 + (2-3)^2 + (3-3.5)^2 + (5-4)^2 + (1-4.5)^2 \\
&= 2^2 + (-1)^2 + (-0.5)^2 + 1^2 + (-3.5)^2 \\
&= 4 + 1 + 0.25 + 1 + 12.25 = 18.5
\end{aligned}$$

Step 3: Calculate SST (Total Sum of Squares)

$$\begin{aligned}
SST &= \sum (y_i - \bar{y})^2 \\
&= (4-3)^2 + (2-3)^2 + (3-3)^2 + (5-3)^2 + (1-3)^2 \\
&= 1^2 + (-1)^2 + 0^2 + 2^2 + (-2)^2 \\
&= 1 + 1 + 0 + 4 + 4 = 10
\end{aligned}$$

Step 4: Calculate R²

$$R^2 = 1 - \frac{SSE}{SST} = 1 - \frac{18.5}{10} = 1 - 1.85 = -0.85$$

Interpretation: Negative R^2 indicates the model performs worse than simply using the mean as prediction. The predictions are very poor.

4.5 Adjusted R-squared

Adjusted R-squared (for multiple regression):

$$R^2_{adj} = 1 - \left(\frac{(1 - R^2)(n - 1)}{n - p - 1}\right)$$

Where:

- n = Number of observations
- p = Number of predictors

Purpose: Penalizes adding unnecessary variables that don't improve model fit.

5 Comparison and Relationships

5.1 Comparison Table

Metric	Purpose	Formula	Range
Margin of Error	Quantifies uncertainty in estimates	$MOE = Z_{\alpha/2} \times \frac{s}{\sqrt{n}}$	$[0, \infty)$
MSE	Measures average squared prediction error	$MSE = \frac{1}{n} \sum (y_i - \hat{y}_i)^2$	$[0, \infty)$
RMSE	MSE in original units	$RMSE = \sqrt{MSE}$	$[0, \infty)$
R²	Proportion of variance explained	$R^2 = 1 - \frac{SSE}{SST}$	$(-\infty, 1]$
Adjusted R²	R² adjusted for number of predictors	$R^2_{adj} = 1 - \frac{(1-R^2)(n-1)}{n-p-1}$	$(-\infty, 1]$

5.2 Mathematical Relationships

- $MSE = \frac{SSE}{n}$
- $R^2 = 1 - \frac{SSE}{SST} = 1 - \frac{n \times MSE}{SST}$
- $RMSE = \sqrt{MSE}$
- Perfect prediction: $MSE = 0$, $R^2 = 1$
- Predicting mean: $\hat{y}_i = \bar{y}$, then $MSE = \frac{SST}{n}$, $R^2 = 0$

5.3 When to Use Each Metric

Theoretical Foundations

Selection Guidelines:

- **Margin of Error:** For statistical estimates and survey results
- **MSE:** For model optimization and training (differentiable)
- **RMSE:** For reporting results in interpretable units
- **R^2:** For explaining model performance to stakeholders
- **Adjusted R^2:** For comparing models with different numbers of predictors

6 Practical Applications in Data Science & ML

Applications in Data Science & Machine Learning

Margin of Error Applications:

- **Survey Design:** Determine sample size for desired precision
- **Clinical Trials:** Report treatment effects with confidence intervals
- **Market Research:** Estimate market share with uncertainty bounds
- **Quality Control:** Monitor process parameters with control limits

Applications in Data Science & Machine Learning

MSE/RMSE Applications:

- **Model Training:** Loss function for regression algorithms
- **Hyperparameter Tuning:** Objective for grid/random search
- **Model Selection:** Compare performance of different algorithms
- **Feature Engineering:** Evaluate impact of new features on error

Applications in Data Science & Machine Learning

R-squared Applications:

- **Business Reporting:** Explain model effectiveness to non-technical stakeholders
- **Model Diagnostics:** Identify underfitting/overfitting
- **Comparative Analysis:** Benchmark against industry standards
- **Academic Research:** Report explanatory power in publications

6.1 Complete Model Evaluation Workflow

1. **Data Preparation:** Split into training (70%), validation (15%), test (15%)

2. **Model Training:** Fit model on training data
3. **Validation:** Evaluate on validation set using MSE/RMSE
4. **Hyperparameter Tuning:** Optimize parameters to minimize validation error
5. **Final Evaluation:** Test on held-out test set
6. **Reporting:** Present MSE, RMSE, R^2 with confidence intervals
7. **Monitoring:** Track metrics over time for model drift

7 Practice Problems

Problem 1: Margin of Error Calculation
A pharmaceutical company wants to estimate the average recovery time for a new treatment. They sample 64 patients and find:

- Sample mean recovery time: 12.5 days
- Sample standard deviation: 3.2 days
- Desired confidence level: 95%

Calculate:

(a) The margin of error

(b) The 95% confidence interval

(c) The required sample size for a margin of error of 0.5 days

Problem 2: MSE and RMSE Calculation
A weather prediction model gives the following forecasts vs actual temperatures (°C):

Actual	**Predicted**
22	20
25	26
18	19
30	28
20	22

Calculate:

(a) The Mean Squared Error (MSE)

(b) The Root Mean Squared Error (RMSE)

(c) If the model improves and all predictions become 1°C closer to actual values, what is the new MSE?

Problem 3: R^2 Interpretation
A housing price prediction model has $R^2 = 0.65$ on the validation set.

(a) What percentage of price variation is explained by the model?

(b) If adding square footage as a feature increases R^2 to 0.72, is this improvement meaningful?

(c) What might cause R^2 to be negative on the test set?

(d) How would you explain this R^2 value to a real estate agent?

Problem 4: Comprehensive Model Evaluation
You're evaluating two regression models for sales prediction:

- Model A: Training MSE = 1500, Validation MSE = 1800, $R^2 = 0.75$
- Model B: Training MSE = 1200, Validation MSE = 1700, $R^2 = 0.78$

Both models were trained on 1000 samples with 10 features.

(a) Calculate RMSE for both models on validation set

(b) Which model shows more overfitting? Why?

(c) Calculate adjusted R^2 for both models

(d) Which model would you deploy and why?

Problem 5: Real-world Scenario
An e-commerce company A/B tests two recommendation algorithms on 10,000 users:

- Algorithm X: Average purchase amount error = \$15 ± \$3 (95% CI)
- Algorithm Y: Average purchase amount error = \$12 ± \$4 (95% CI)

(a) Which algorithm has better average performance?

(b) Which algorithm has more precise estimates?

(c) Calculate the confidence intervals for both algorithms

(d) Design a statistical test to determine if the difference is significant

8 Implementation Guidelines

8.1 Python Implementation

```
import numpy as np
from sklearn.metrics import mean_squared_error, r2_score
from scipy import stats

def calculate_metrics(y_true, y_pred, confidence=0.95):
    """Calculate comprehensive error metrics"""

    # Basic metrics
    n = len(y_true)
    mse = mean_squared_error(y_true, y_pred)
    rmse = np.sqrt(mse)
    r2 = r2_score(y_true, y_pred)

    # Confidence interval for MSE
    errors = y_true - y_pred
    se_mse = np.std(errors**2) / np.sqrt(n)
    z_score = stats.norm.ppf(1 - (1-confidence)/2)
    mse_ci = (mse - z_score*se_mse, mse + z_score*se_mse)
```

```
    # Adjusted R-squared
    p = 1  # Assuming 1 predictor for simple regression
    r2_adj = 1 - ((1 - r2) * (n - 1) / (n - p - 1))

    return {
        'MSE': mse,
        'MSE_CI': mse_ci,
        'RMSE': rmse,
        'R2': r2,
        'R2_adj': r2_adj,
        'Sample_Size': n
    }

# Example usage
y_actual = np.array([4, 2, 3, 5, 1])
y_pred = np.array([2, 3, 3.5, 4, 4.5])
metrics = calculate_metrics(y_actual, y_pred)
print(metrics)
```

8.2 R Implementation

```
# Error metrics calculation in R
calculate_metrics <- function(y_true, y_pred, confidence = 0.95) {
  n <- length(y_true)
  errors <- y_true - y_pred

  # Basic metrics
  mse <- mean(errors^2)
  rmse <- sqrt(mse)
  sse <- sum(errors^2)
  sst <- sum((y_true - mean(y_true))^2)
  r2 <- 1 - (sse / sst)

  # Confidence interval for MSE
  se_mse <- sd(errors^2) / sqrt(n)
  z_score <- qnorm(1 - (1 - confidence) / 2)
  mse_ci <- c(mse - z_score * se_mse, mse + z_score * se_mse)

  # Adjusted R-squared
  p <- 1  # Number of predictors
  r2_adj <- 1 - ((1 - r2) * (n - 1) / (n - p - 1))

  list(
    MSE = mse,
    MSE_CI = mse_ci,
    RMSE = rmse,
    R2 = r2,
    R2_adj = r2_adj,
    Sample_Size = n
  )
}
```

```
# Example usage
y_actual <- c(4, 2, 3, 5, 1)
y_pred <- c(2, 3, 3.5, 4, 4.5)
metrics <- calculate_metrics(y_actual, y_pred)
print(metrics)
```

8.3 Best Practices

Theoretical Foundations

Implementation Guidelines:

- **Always use multiple metrics:** No single metric tells the whole story
- **Report confidence intervals:** Especially for MSE and error rates
- **Use appropriate baselines:** Compare against simple models (mean, median)
- **Consider business context:** Different applications prioritize different metrics
- **Document assumptions:** Note data preprocessing and transformations
- **Monitor over time:** Track metrics for model drift detection

9 Summary and Key Takeaways

9.1 Metric Selection Guide

Use Case	Primary Metric	Secondary Metric
Statistical Estimation	Margin of Error	Confidence Interval Width
Model Optimization	MSE (differentiable)	Validation RMSE
Model Comparison	Adjusted R^2	MSE with statistical test
Business Reporting	R^2 (interpretable)	RMSE (practical units)
Algorithm Selection	Cross-validated MSE	R^2 with confidence intervals

9.2 Common Pitfalls and Solutions

Learning Resources

Pitfalls to Avoid:

- **Overfitting to R^2:** High R^2 doesn't guarantee good predictions on new data
- **Ignoring confidence intervals:** Point estimates can be misleading
- **Data leakage:** Ensure proper train/validation/test splits
- **Scale dependence:** MSE is sensitive to data scale; consider normalization
- **Multiple comparisons:** Adjust significance levels when comparing many models
- **Assuming normality:** Check error distribution assumptions

9.3 Future Learning Path

Book Chapter Reference

Recommended Resources:

- **Books:**
  - "Introduction to Statistical Learning" by James, Witten, Hastie, Tibshirani
  - "Elements of Statistical Learning" by Hastie, Tibshirani, Friedman
  - "Statistical Rethinking" by Richard McElreath
- **Online Courses:**
  - Coursera: Machine Learning by Andrew Ng
  - edX: Statistics and Data Science MicroMasters
  - Kaggle: Machine Learning courses
- **Practical Tools:**
  - Python: scikit-learn, statsmodels, scipy
  - R: caret, mlr, tidymodels
  - Visualization: matplotlib, seaborn, ggplot2

9.4 Reference Table: Critical Values

Confidence Level	Z-value	t-value (n=30)	t-value (n=100)
90%	1.645	1.699	1.660
95%	1.960	2.045	1.984
99%	2.576	2.756	2.626

"The most important questions in life are, for the most part, only problems in probability." — Pierre-Simon Laplace

"In God we trust; all others must bring data." — W. Edwards Deming

---